

Operating Kubernetes on IBM Cloud

Part 1 - Essentials

—
Aco Vidović
Hybrid Cloud Build Team Leader
IBM Central & Eastern Europe



Agenda

- Part 1 - Lecture:
 - Kubernetes objects
 - Working with Kubernetes Command Line Interface (CLI)
 - Application autoscaling and rolling updates
 - Deploying Kubernetes to IBM Cloud

Break

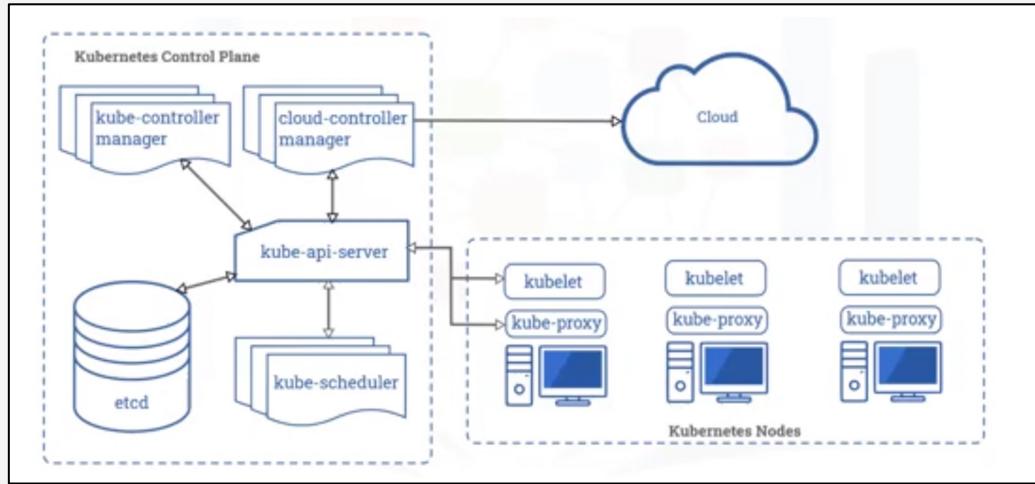


- Part 2 – Demo and hands-on labs:
 - Create your free instance of IKS
 - Deploy container image to Kubernetes
 - Updating apps with rolling updates and rollbacks
 - Using YAML configuration files
 - Scaling and auto-scaling
 - Connecting to a back-end service

Access materials:

<https://ibm.biz/operating-iks-workshop>

Managed Kubernetes cluster

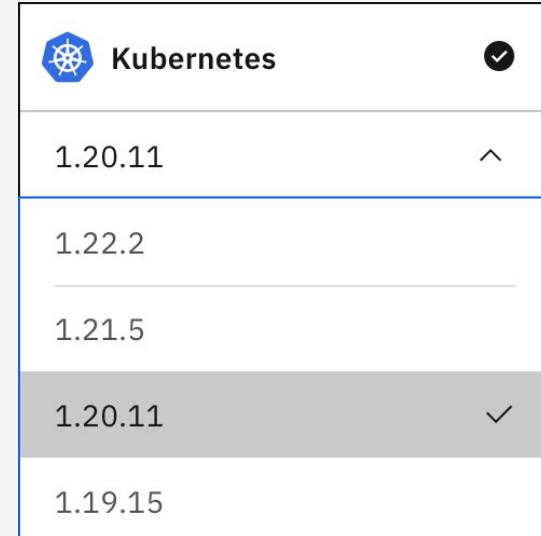


- Kubernetes cluster is made of:
 - One or more **Nodes** (also known as Worker nodes). This is where apps run.
 - **Control Plane** (also known as the Master node) which manages nodes
- Cloud Controller Manager in the Control Plane talks to the underlying cloud infrastructure
- Nodes are created by the cloud provider, not by Kubernetes

[More info...](#)

IBM Kubernetes Service

- A **certified, managed Kubernetes service** providing an intuitive user experience with on-going cluster management.
- Built-in **security and isolation** to enable rapid delivery of apps.
- Available in six IBM regions WW, including **20+ datacenters**.
- Fully dedicated, **single tenant Kubernetes clusters** deployed within customer account and network
- Seamless **integrated with IBM Cloud services**
- **Portability** with native Kubernetes experience and full **API support**



[More info on IKS...](#)

Our hands-on exercises today

There are 2 ways to run hands-on lab today, choose the one you prefer:

- Use a demo environment at [IBM Developer web site](#). No setup needed.
- Run the lab on your own Kubernetes cluster:
 - You create a cluster run by the IBM Kubernetes Service (IKS)
 - Download lab instructions files to your computer
 - Use CLI on your computer to interact with your IKS cluster

Preparing your own lab environment (optional)

If you choose to run the lab in your own environment:

1. Start a creation of a free Kubernetes cluster
 - In your web browser, go to <https://ibm.biz/operating-iks-workshop>
 - Download the instructions in file *Lab-1-create-iks-cluster.pdf*
 - Read and follow the instructions in the PDF file
2. Allow the cluster creation process enough time to finish. In the meantime, we will continue with lecture.
3. After 15-20 min, come back to the lab to complete it by setting up all necessary CLI tools.

[Demo of this lab](#)

Kubernetes objects

Kubernetes objects

- Kubernetes objects are **persistent entities**. They represent the state of your cluster and define the desired state for your workload.
- Work with objects using Kubernetes APIs through:
 - kubectl CLI or
 - your programs using one of the [Client Libraries](#)
- Kubernetes has many **types** of objects: Nodes, Pods, Deployments, Services, Volumes....
 - Containers are not Kubernetes objects, they live in Pods
- Each object has configuration that includes (among others) two parts: **spec** and **status**
 - *Spec* depicts the desired state of object – as you want it to be
 - **A human being (you)** usually defines spec
 - *Status* of the object is what currently is
 - Kubernetes makes sure the status matches your spec
 - Kubernetes keeps track of the status. You can see it if you want to.

Node

- A type of Kubernetes object
- The worker machine in Kubernetes
 - Can be virtual or physical machine
- Managed by the Kubernetes control plane



Container runtime in Node

- Downloads images from registries and runs containers
- Kubernetes implements an interface so that this component is pluggable
- Available runtimes:
 - Docker – the most well known
 - CRI-O
 - containerd
 - and any other implementation of the [Kubernetes Container Runtime Interface \(CRI\)](#)

[More info...](#)

Labels and selectors

Labels

- Kubernetes objects have them. Purpose:
 - Identification of objects
 - Organization and grouping of objects
- Key/value pairs **attached to objects**
- Labels are not unique
 - Same label can be assigned to many objects.
 - Example: give the same label to all objects belonging to the same app
 - Each object can have zero, one or more labels
- Examples of labels:
 - “**app**” : “**ticketing**” – marks an object belonging to an *app* named “*ticketing*”
 - “**stage**” : “**prod**” – marks an object belonging to the production stage

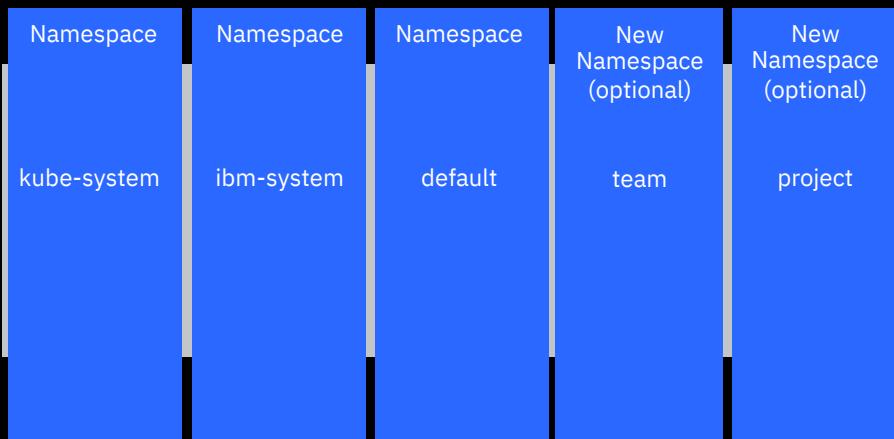


Selectors

- Kubernetes’ means to select objects using their labels

Namespace

- Namespace is a type of **object**
- It helps with virtualisation of a physical cluster
- Segregate cluster by some meaningful criteria such as team, project etc.
- Necessary with larger number of users



Object names

- Each object has a name
- Names are **unique** for a type of resource **within a namespace**. You can give the same name to:
 - Two objects of *different types* in the same namespace:
 - A pod and a service named *guestbook* in namespace *default*
 - Two objects of the same type in *different namespaces*:
 - Pod named *guestbook* in namespace *default* and a pod *guestbook* in namespace *team*

Pod

- Pod is a smallest unit of work in Kubernetes
- It represents processes running in your cluster
- Container is in a Pod
 - Rarely more than one container in a pod
 - Replicating a pod scales an application horizontally
- Like other Kubernetes objects, Pod has a configuration described in its configuration file:
 - Either in YAML format
 - Or JSON

These four fields are present in every Kubernetes configuration file

YAML configuration file describing a Pod

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   labels:
      app: guestbook
5 spec:
6   containers:
7     - image: ibmcom/guestbook:v1
      name: guestbook
      ports:
        - containerPort: 3000
```

Deployment

- An object that provides **updates** to Pod
- Contains a Pod configuration
 - In the template field
- Can run multiple **replicas** of a pod
 - Replica is identical copy of a pod
 - For scalability and redundancy
- Suitable for stateless applications
- A Deployment update triggers a Pod rollout
- Deployment controls:
 - Rollout of a pod to new version
 - Roll back to previous version

YAML file describing a Deployment

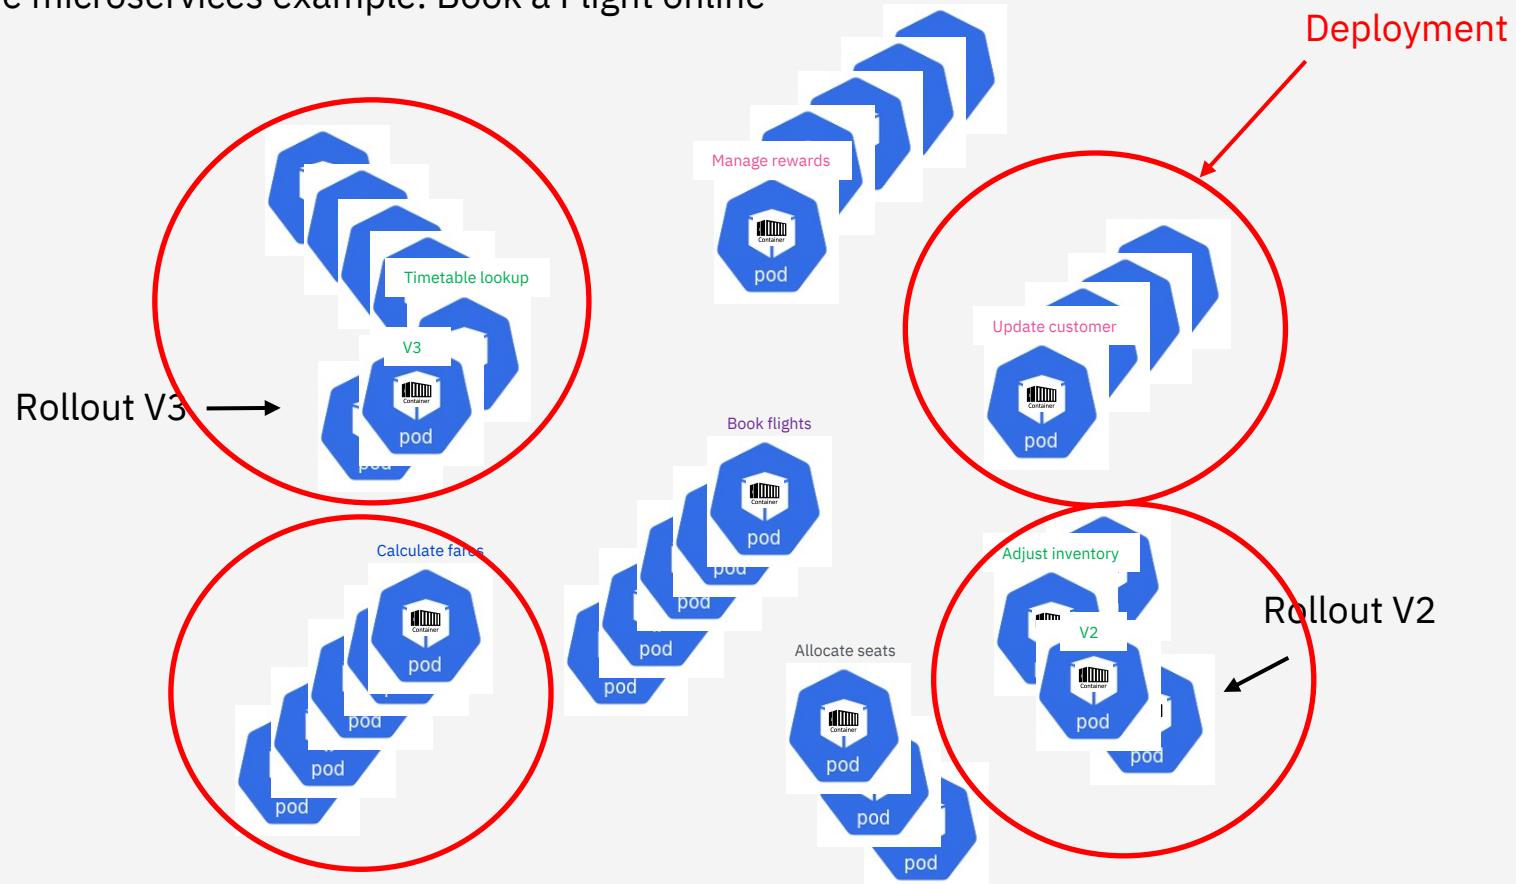
```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: guestbook-v1
  labels:
    app: guestbook
    version: "1.0"
spec:
  replicas: 3
  selector:
    matchLabels:
      app: guestbook
  template:
    metadata:
      labels:
        app: guestbook
        version: "1.0"
    spec:
      containers:
        - name: guestbook
          image: ibmcom/guestbook:v1
          ports:
            - name: http-server
              containerPort: 3000
```

Pod definition



Why do we need Deployments

A real-life microservices example: Book a Flight online



ReplicaSet

- ReplicaSet is a type of object
- Like Deployment, it maintains a set of identical pods
- Configuration looks similar to Deployment. It has:
 - Number of replicas
 - Pod template
 - Selector to identify which Pods it can acquire
- When you create a Deployment, ReplicaSet is automatically created
- It is possible but not recommended to explicitly create a ReplicaSet. Better create a Deployment, because Deployment:
 - Does everything that ReplicaSet does
 - Plus it can do application rollout and roll back

YAML file describing a ReplicaSet

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  labels:
    app: guestbook
spec:
  replicas: 3
  selector:
    matchLabels:
      app: guestbook
  template:
    metadata:
      labels:
        app: guestbook
    spec:
      containers:
        - image: ibmcom/guestbook:v1
          name: guestbook
        ports:
          - containerPort: 3000
```

Service

- Service is a type of Kubernetes object
- It makes Deployments and Pods available to the outside world
- Your users or other apps can not directly access Deployments or Pods
 - Deployment/Pod **must be exposed through a Service**, only then your users and apps can access them
 - Pods are not-permanent, they are created and destroyed, while **Services are permanent**
- Service types:
 - ClusterIP – only available from within a cluster
 - NodePort - available via *IPaddress:Port*
 - LoadBalancer – available via a Cloud provider's load balancer
 - ExternalName – available via a CNAME record

YAML file describing a Service

```
apiVersion: v1
kind: Service
metadata:
  name: guestbook
  labels:
    app: guestbook
spec:
  ports:
  - port: 3000
    targetPort: http-server
  selector:
    app: guestbook
  type: LoadBalancer
```

Working with IKS via Command Line Interface (CLI)

IBM Cloud CLI

- The CLI to work with IBM Cloud and all its resources
- The command is: `ibmcloud`
- It has several plugins, including:
 - IBM Kubernetes service (IKS). Command: `ibmcloud ks`
 - Container registry. Command: `ibmcloud cr`

Kubernetes CLI - `kubectl`

- Key tool for working with Kubernetes – platform independent
- Provides functionality for working with Kubernetes clusters and their workloads
- Two command types:
 - Imperative commands:
 - quickly create, update, and delete Kubernetes objects
 - easiest to learn
 - Declarative commands:
 - configuration files define one or more objects
 - no operation is specified

Working with IKS from CLI – sequence of use

1. Establish connection to your IBM cloud account.

In terminal window on your computer, run command:

```
ibmcloud login -a cloud.ibm.com -r <region> -g <resource-group>
```

2. Establish connection (set the Kubernetes context) to your IKS cluster.

On your computer run:

```
ibmcloud ks cluster config --cluster <your-cluster-name>
```

3. Verify that you can connect to your IKS cluster.

On your computer run:

```
kubectl config current-context
```

Kubernetes imperative commands examples

```
kubectl run nginx -image nginx
```

Pros of imperative commands:

- Quickly create, update, and delete Kubernetes objects
- Easiest to learn

Cons

- Don't provide any audit trail
- Not very flexible

```
kubectl create -f nginx.yaml
```

- Use configuration template
- Specify an operation such as create, replace, or delete

Declarative commands

- Configuration files define one or more objects
- No operation is specified
 - Relevant operations are inferred by kubectl
- Work on files and directories
- Configuration files define desired state, and Kubernetes actualizes that state
- Preferred method for production systems

```
kubectl apply -f nginx.yaml
```

Viewing Kubernetes resources with *kubectl* command

Two popular commands to view your Kubernetes resources:

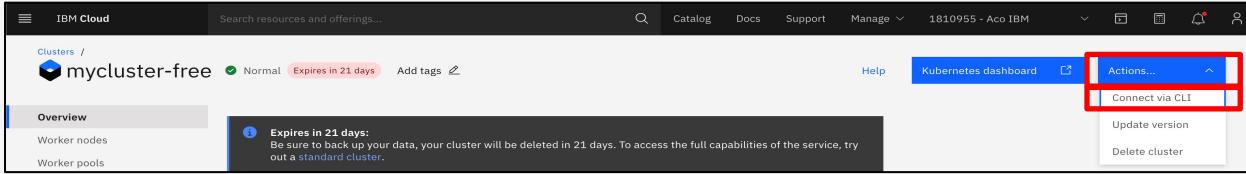
- `kubectl get`: to list resources and get basic data about them
- `kubectl describe`: to get details for a specific resource

```
kubectl get deployments -namespace kube-system
```

```
kubectl describe deployment kube-dns-amd64 -namespace kube-system
```

Finish the set up of your Kubernetes cluster (optional)

1. In your IBM Cloud console, go to your cluster, click **Actions**, then **Connect via CLI**



2. Follow the instructions in pop-up window

Connect via CLI

Cluster status: Normal

If this is your first time connecting to an IBM Cloud cluster, see the [full setup directions](#).

1. Log in to your IBM Cloud account. Include the `--ssso` option if using a federated ID.

```
ibmcloud login -a cloud.ibm.com -r eu-de -g default
```

2. Set the Kubernetes context to your cluster for this terminal session. For more information about this command, [see the docs](#).

```
ibmcloud ks cluster config --cluster c5mnod2f0vl6ngs4p86g
```

3. Verify that you can connect to your cluster.

```
kubectl config current-context
```

Now, you can run `kubectl` commands to manage your cluster workloads in IBM Cloud! For a full list of commands, see the [Kubernetes docs](#).

Tip: Plan to use multiple clusters? Repeat these steps for each cluster. Then, you can use the `kubectl config use-context` command to switch your context to a different cluster.

Often used commands: listing resources with `kubectl get`

```
😊 => kubectl get deployments
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
guestbook-v1	3/3	3	3	26h
redis-master	1/1	1	1	26h
redis-slave	2/2	2	2	26h

```
😊 => kubectl get namespaces
```

NAME	STATUS	AGE
default	Active	28h
ibm-cert-store	Active	28h
ibm-operators	Active	28h
ibm-system	Active	28h
kube-node-lease	Active	28h
kube-public	Active	28h
kube-system	Active	28h

```
😊 => kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
guestbook-v1-7676687667-klw2q	1/1	Running	0	26h
guestbook-v1-7676687667-s2gw1	1/1	Running	0	26h
guestbook-v1-7676687667-wvfv4m	1/1	Running	0	26h
redis-master-7bf5f6d487-t9t89	1/1	Running	0	26h
redis-slave-58db457857-5pngx	1/1	Running	0	26h
redis-slave-58db457857-shgh9	1/1	Running	0	26h

```
😊 => kubectl get services
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
guestbook	LoadBalancer	172.21.80.230	<pending>	3000:30614/TCP	27h
kubernetes	ClusterIP	172.21.0.1	<none>	443/TCP	28h
redis-master	ClusterIP	172.21.183.178	<none>	6379/TCP	26h
redis-slave	ClusterIP	172.21.254.149	<none>	6379/TCP	26h

```
😊 => kubectl get services -o wide
```

NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE	SELECTOR
guestbook	LoadBalancer	172.21.70.225	<pending>	3000:32338/TCP	17m	app=guestbook
kubernetes	ClusterIP	172.21.0.1	<none>	443/TCP	29h	<none>
redis-master	ClusterIP	172.21.183.178	<none>	6379/TCP	27h	app=redis,role=master
redis-slave	ClusterIP	172.21.254.149	<none>	6379/TCP	27h	app=redis,role=slave

There is a long / short command version :

- `kubectl get pods` OR `kubectl get po`
- `deployments` OR `deploy`
 - `services` OR `svc`
 - `namespaces` OR `ns`
 - etc...

Help: `kubectl get -h`

For more output:
`-o wide`

Obtain details about any object with `kubectl describe`

```
[☺ => kubectl describe deployment guestbook-v1
Name:                      guestbook-v1
Namespace:                 default
CreationTimestamp:         Sun, 28 Feb 2021 19:07:40 +0100
Labels:                    app=guestbook
                           version=1.0
Annotations:               deployment.kubernetes.io/revision: 1
Selector:                  app=guestbook
Replicas:                  3 desired | 3 updated | 3 total | 3 available | 0 unavailable
StrategyType:              RollingUpdate
MinReadySeconds:           0
RollingUpdateStrategy:    25% max unavailable, 25% max surge
Pod Template:
  Labels:  app=guestbook
          version=1.0
  Containers:
    guestbook:
      Image:      ibmcom/guestbook:v1
      Port:       3000/TCP
      Host Port:  0/TCP
      Environment: <none>
      Mounts:     <none>
      Volumes:    <none>
  Conditions:
    Type        Status  Reason
    ----        ----   -----
    Available   True    MinimumReplicasAvailable
    Progressing True    NewReplicaSetAvailable
  OldReplicaSets: <none>
  NewReplicaSet:  guestbook-v1-7676687667 (3/3 replicas created)
  Events:       <none>
```

`kubectl describe deployment`

`kubectl describe svc`

`kubectl describe pod`

Etc...

Creating and deleting resources with *kubectl*

Create object imperatively:

- *kubectl create <resource-type> <resource-name> ...*
- *kubectl create -f <file-name>*

Declaratively: *kubectl apply -f <file-name>*

Delete: *kubectl delete <resource-type> <resource-name>*

```
😊 => kubectl apply -f guestbook-service.yaml
service/guestbook created
```

```
😊 => kubectl get service guestbook
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
guestbook  LoadBalancer  172.21.70.225  <pending>      3000:32338/TCP  49s
```

```
[😊 => kubectl delete service guestbook
service "guestbook" deleted
```

Manage Kubernetes applications

Some common operations. Let Kubernetes do:

- **Auto-scaling** – automatically scale your app up or down
- Rolling updates – update applications to another version while they are still running

Auto-scaling – how to do it?

You can auto-scale a Deployment or a ReplicaSet using command `kubectl autoscale`

Example: auto-scale between 2 and 4 replicas while keeping the CPU utilization around 10%:

```
😊 => kubectl autoscale deployment guestbook-v1 --min=2 --max=4 --cpu-percent=10
horizontalpodautoscaler.autoscaling/guestbook-v1 autoscaled
```

```
😊 => kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
guestbook-v1-7676687667-c2dkf  0/1     Terminating   0   11m
guestbook-v1-7676687667-c66mp  1/1     Running     0   28m
guestbook-v1-7676687667-m2ghg  1/1     Running     0   16m
guestbook-v1-7676687667-s2jxd  0/1     Terminating   0   16m
guestbook-v1-7676687667-vtxwv  1/1     Running     0   28m
guestbook-v1-7676687667-zqzkq  1/1     Running     0   16m
```

And after a while...

```
😊 => kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
guestbook-v1-7676687667-c66mp  1/1     Running   0   30m
guestbook-v1-7676687667-m2ghg  1/1     Running   0   18m
guestbook-v1-7676687667-vtxwv  1/1     Running   0   30m
guestbook-v1-7676687667-zqzkq  1/1     Running   0   18m
redis-master-7bf5f6d487-t9t89  1/1     Running   0   29h
```

Auto-scaling – behind the scenes

A Horizontal Pod Autoscaler (HPA) was created behind the scenes

[😊 => kubectl get hpa

NAME	REFERENCE	TARGETS	MINPODS	MAXPODS	REPLICAS	AGE
guestbook-v1	Deployment/guestbook-v1	<unknown>/10%	2	4	4	9m19s

You can create a standalone HPA using YAML,
but it is recommended autoscale with *kubectl*

```
apiVersion: autoscaling/v1
kind: HorizontalPodAutoscaler
metadata:
  name: guestbook-v1
  namespace: default
spec:
  maxReplicas: 4
  minReplicas: 2
  scaleTargetRef:
    apiVersion: apps/v1
    kind: Deployment
    name: guestbook-v1
  targetCPUUtilizationPercentage: 10
```

Manage Kubernetes applications

Some common operations.

- Auto-scaling
- Rolling updates – update applications to another version while they are still running

Application update strategies

- **Recreate:** *terminate the old app version, then release new one*
- **Rolling:** at the same time terminate old pods and create new ones, so some pods are always running. *There is no app downtime*
- **Blue-green:** release a new version alongside the old version then *switch traffic*
- **Canary:** release a *new version to a subset of users*, then proceed to a full rollout

Rolling updates

- Replicas and autoscaling are important to *minimize downtime* and service interruptions
- Rolling updates are a way to roll out application changes in an automated and controlled fashion throughout your pods
- Work with Pod templates defined in Deployments
- Allow for rollback if something went wrong with new version

Enable your Deployment configuration for rolling updates

1. Add *liveness* and *readiness probes* to your configuration file. This ensures deployments are **marked ready** appropriately.

Add these two sections into the Deployment's config file:

```
livenessProbe:  
  httpGet:  
    path: /  
    port: 9080  
  initialDelaySeconds: 300  
  periodSeconds: 15  
readinessProbe:  
  httpGet:  
    path: /  
    port: 9080  
  initialDelaySeconds: 45  
  periodSeconds: 5
```

2. Add rolling update strategy to your YAML file.

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: nginx-test  
spec:  
  replicas: 10  
  selector:  
    matchLabels:  
      service: http-server  
  minReadySeconds: 5  
  progressDeadlineSeconds: 600  
strategy:  
  type: RollingUpdate  
  rollingUpdate:  
    maxUnavailable: 50%  
    maxSurge: 2
```

Rolling updates: rollout the new version - V2

```
kubectl set image...
```

```
[😊 => kubectl set image deployment/guestbook-v1 guestbook=ibmcom/guestbook:v2  
deployment.apps/guestbook-v1 image updated
```

To check what happened: `kubectl rollout status...`

```
[😊 => kubectl rollout status deployments/guestbook-v1  
deployment "guestbook-v1" successfully rolled out
```

```
[😊 => kubectl describe po guestbook-v1-8489ddbb68-2bpfn  
Name:           guestbook-v1-8489ddbb68-2bpfn  
Namespace:      default  
Priority:      0  
Node:          10.144.51.71/10.144.51.71  
Start Time:    Tue, 02 Mar 2021 01:06:25 +0100  
Labels:         app=guestbook  
               pod-template-hash=8489ddbb68  
               version=1.0  
Annotations:   cni.projectcalico.org/podIP: 172.30.97.22/32  
               cni.projectcalico.org/podIPs: 172.30.97.22/32  
               kubernetes.io/psp: ibm-privileged-psp  
Status:        Running  
IP:            172.30.97.22  
IPs:  
  IP:          172.30.97.22  
Controlled By: ReplicaSet/guestbook-v1-8489ddbb68  
Containers:  
  guestbook:  
    Container ID:  containerd://b5943c6de620b9371bc53c1c2e263021ee845d85c7be518bf76c1259791326f0  
    Image:         ibmcom/guestbook:v2
```

Rolling updates: rollback to previous version

If the update doesn't work well: `kubectl rollout undo...`

```
[😊 => kubectl rollout undo deployment/guestbook-v1
deployment.apps/guestbook-v1 rolled back
```

```
[😊 => kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
guestbook-v1-7676687667-6wqqq	0/1	ContainerCreating	0	1s
guestbook-v1-7676687667-gzwmg	1/1	Running	0	4s
guestbook-v1-7676687667-mx768	0/1	ContainerCreating	0	0s
guestbook-v1-7676687667-thtwj	1/1	Running	0	4s
guestbook-v1-8489ddbb68-2bpfn	1/1	Running	0	7m25s
guestbook-v1-8489ddbb68-frzgz	1/1	Terminating	0	7m22s
guestbook-v1-8489ddbb68-qntt2	1/1	Terminating	0	7m25s
guestbook-v1-8489ddbb68-vv8px	0/1	Terminating	0	7m22s

Rolling updates: check the rollback status

```
[😊 => kubectl rollout status deployments/guestbook-v1
deployment "guestbook-v1" successfully rolled out
```

```
[😊 => kubectl describe po guestbook-v1-7676687667-6wqqq
Name:           guestbook-v1-7676687667-6wqqq
Namespace:      default
Priority:       0
Node:          10.144.51.71/10.144.51.71
Start Time:    Tue, 02 Mar 2021 01:13:49 +0100
Labels:         app=guestbook
                pod-template-hash=7676687667
                version=1.0
Annotations:   cni.projectcalico.org/podIP: 172.30.97.18/32
                cni.projectcalico.org/podIPs: 172.30.97.18/32
                kubernetes.io/psp: ibm-privileged-psp
Status:        Running
IP:            172.30.97.18
IPs:
  IP:          172.30.97.18
Controlled By: ReplicaSet/guestbook-v1-7676687667
Containers:
  guestbook:
    Container ID:  containerd://f3dd26af8204b2f5f485e4c3b182b40cb2c31298dcf904b16bab132d1b379234
    Image:         ibmcom/guestbook:v1
```

Hands-on lab

In this lab, you will work with Deployments, Services, scaling and updates on IKS

- Deploy container image with the Guestbook app to your cluster
- Experiment with replicas, application update and rollback
- Use YAML configuration files
- Persist data by connecting the app to a back-end service (database)

How to do the hands-on lab

Option 1: Do the lab using your own Kubernetes cluster on IBM Cloud

1. In your web browser, go to <https://ibm.biz/operating-iks-workshop>
2. Download the file *Lab-2-iks-essentials.pdf*
3. Read and follow the instructions in the PDF file

Option 2: Do the lab in a pre-configured demo environment (faster)

- In a demo environment at [IBM Developer web site](https://developer.ibm.com/openlabs/cloudlabs/containers-and-kubernetes-essentials), click on **Launch Lab** and follow instructions.



Resources

Part 2 of this workshop: Operating Kubernetes on IBM Cloud, Part 2 - Best Practices

- November 3 at 16:00 – 18:00 CET
- Same URL as today

[IBM Kubernetes Service docs](#)

Videos: Introduction to Container Orchestration with IBM Cloud, [Part 1](#) and [Part 2](#)

[Credly \(Acclaim\) badge - Operating Kubernetes on IBM Cloud](#)

Thank you!